

AI-based Meeting Notes Generator Using NLP Techniques

1st Siddhi S Kakani
Artificial Intelligence and Data Science
Boston, USA
kakani.s@northeastern.edu

2nd Simona Matiukaite
Computer Science
Boston, USA
matiukaite.s@northeastern.edu

Abstract—As organizations worldwide continue shifting toward remote collaboration, the number of virtual meetings held daily across corporate, academic, and clinical environments has grown substantially. Manual note-taking during these meetings is often inefficient and error-prone, resulting in missed action items, overlooked deadlines, and forgotten decisions. While several commercial platforms offer meeting summarization capabilities, most are limited in input modality, lack domain-specific specialization, and provide little transparency regarding their underlying models. This paper presents an AI-based Meeting Decision and Action Intelligence System built to overcome these shortcomings. The architecture follows a sequential processing chain that ingests video, audio, or textual meeting content and generates structured outputs including meeting summaries, speaker-assigned action items, deadline extraction, key decision identification, and domain classification. The pipeline employs the Whisper speech recognition model developed by OpenAI alongside two generative summarization architectures — BART and FLAN-T5 — for abstractive summarization, and pattern-based extraction for actionable information. A hybrid summarization strategy combining abstractive and extractive approaches is employed to maximize both content coverage and coherence. The models are fine-tuned on the AMI Meeting Corpus and evaluated across three application domains using standard ROUGE metrics. The system is deployed as a web-based interface using Gradio, enabling real-time meeting analysis from multiple input modalities. By combining transparent model selection with multi-domain evaluation, this work contributes a reproducible and extensible framework for automated meeting intelligence.

Index Terms — BART, FLAN-T5, Whisper, Abstractive Text Summarization, Meeting Intelligence, Speech-to-Text, Gradio.

I. INTRODUCTION

In recent decades, rapid technological advancement has substantially accelerated the pace of daily life and work. At the same time, the widespread adoption of hybrid and fully remote work environments has increased individuals' exposure to distractions outside the immediate work context, thereby raising the likelihood of missed details and human error during virtual interactions. A similar issue arises in recorded classroom lectures, which may be viewed as a specialized form of meeting in which important information such as explanations, questions, assignments, and deadlines is communicated. In both professional and educational settings, manual note-taking during online sessions can therefore be inefficient and unreliable, creating a need for technological support. One possible solution is AI-based video summarization.

A variety of online platforms currently offer summarization services for audio, video, or text content. Popular examples include Recal, NoteGPT, Read AI, Jumpshare, MindGrasp,

and GoTranscript. These tools typically take video, audio, or textual input and generate a textual summary as output. However, existing summarization services differ in both input modality and domain emphasis.

First, some platforms are limited in the types of input they support. For example, Read AI [10], Recal [8], and GoTranscript [9] primarily focus on audio and video summarization, which may require users to rely on separate tools when summarizing documents or other text-based material. Second, current tools tend to be oriented toward specific application domains rather than demonstrating broad, clearly defined domain specialization. Read AI, Jumpshare [7], and Recal are largely positioned for business and meeting productivity, whereas MindGrasp [6] and NoteGPT are more closely associated with educational use cases. Although these platforms may be capable of processing content from healthcare or medical contexts, there is limited evidence that they are explicitly designed for clinical summarization, medical compliance, or terminology-heavy healthcare documentation. The ability to summarize healthcare-related content should not be treated as equivalent to domain-specific specialization.

In addition to summarization, several of these tools also provide supplementary features such as action-item extraction, assignee identification, and deadline detection. This is the case for tools such as Recal, NoteGPT, and Read AI. Such features may be useful in both academic and workplace settings, where users need not only condensed content but also actionable task information.

A more significant limitation, however, is the limited transparency surrounding the underlying models used by many of these platforms. In many cases, the providers do not clearly disclose which models are used for transcription, summarization, or action-item extraction. As a result, the performance characteristics and limitations of these systems remain difficult to assess. This raises concerns about the reliability of generated summaries, particularly for longer inputs, multi-speaker discussions, or contexts in which preserving nuance and detail is critical. The lack of transparency also makes it more difficult to judge whether such tools are appropriate for sustained use in domains such as education, business, or healthcare.

These limitations, and especially the lack of transparency regarding the models used in existing summarization tools, motivate this project. This work investigates the effectiveness of two transformer-based summarization models, FLAN-T5-

base and BART-samsum, for generating meeting summaries from content derived from video, audio, and text inputs across business, education, and healthcare domains. The main contribution of this work is a hybrid meeting summarization pipeline that combines Whisper [22] transcription, extractive sentence selection, abstractive summarization, and structured information extraction across multiple domains.

II. LITERATURE REVIEW

The paper “AI-based Meeting Notes Generator using Whisper and BERT” by Kamalesh S., Manojkumar S., and Kavinesh K. K., published at the 2025 ICICNIS conference, introduces a fully automated, AI-driven meeting notes generation system that converts raw audio files into structured meeting minutes, action items, and key discussion points [1]. The system was designed to increase participants’ attention during meetings by eliminating the need for manual note-taking, thereby improving accuracy, enhancing team productivity, and reducing human error that could lead to critical misunderstandings.

In the article, the proposed methodology consists of a modular, sequential pipeline comprising six steps: audio input, speech-to-text conversion, speaker diarization, text segmentation, summarization, and web interface output. The process begins when a user uploads a meeting recording to a web-based interface implemented in React.js. The file is then sent to OpenAI’s Whisper API for transcription. Whisper is a transformer-based automatic speech recognition (ASR) model known for its robustness across different audio qualities, accents, and noise conditions. Following transcription, speaker diarization is performed using pyAnnote, which distinguishes between speakers by labeling segments of speech. Each transcript segment is tagged with the corresponding speaker and timestamp. Once the transcript is prepared, the segmented text is processed through a BERT-based NLP summarization model. BERT (Bidirectional Encoder Representations from Transformers) analyzes contextual relationships within the transcript and identifies key action items and decisions. The final notes are displayed on the web interface and can be downloaded in PDF or DOCX format. The backend is implemented using Python frameworks such as Flask or FastAPI, with data stored in PostgreSQL or MongoDB databases.

All components of the pipeline—including ASR with Whisper, summarization, and action-item extraction—demonstrated consistent performance. The speech recognition module achieved a word error rate (WER) of approximately 12%, indicating high transcription quality. Summarization performance ranged from 0.70 to 0.78 for ROUGE-1 and ROUGE-L scores, respectively, reflecting strong agreement with human-written references. Action-item extraction achieved an F1 score of 0.82, while precision, recall, and F1 scores for individual modules ranged between 0.85 and 0.90. Compared to TextRank and keyword-based approaches, the proposed system improved ROUGE-L by 12% and F1 by 15%, while maintaining an

average processing time of under four seconds per minute of audio.

This article is highly relevant to our project because it shares the goal of developing an AI-based system for extracting meeting decisions and action items from audio recordings. Similar to the authors, we use OpenAI’s Whisper API for speech-to-text transcription and pyAnnote for speaker diarization, resulting in a comparable overall pipeline. However, our project differs in two major ways. First, rather than relying on a single BERT-based summarization model, we compare two transformer-based abstractive summarization models: BART-samsum and FLAN-T5-base. While BERT is primarily used for extractive summarization by selecting important sentences from the original text, both BART and FLAN-T5 generate new sentences that better capture the overall meaning of the discussion. This allows the system to produce more coherent and readable summaries. In addition, comparing two different abstractive models provides a more meaningful evaluation of which architecture is better suited for meeting summarization tasks. Second, unlike the authors who relied primarily on benchmark datasets, our project also includes a domain-specific collection of business, education, and healthcare meetings. By incorporating multiple meeting domains, we aim to evaluate how well the models generalize across different conversational settings. This article therefore provides a useful baseline for comparison and supports several of the architectural choices made in our system.

BERT is only one of many pre-trained models that can be applied to text summarization. A more meaningful research question is whether any single model consistently outperforms other transformer-based approaches, and if so, under what conditions. Since this project compares BART-samsum and FLAN-T5-base, understanding the relative strengths of different summarization architectures is especially important. The article “Exploring Abstractive Summarization of Pre-Trained Models: A Study on GPT-2, T5, PEGASUS and BART” investigates this question in the context of abstractive summarization [2].

In that study, the authors evaluated four pre-trained models—GPT-2, T5, PEGASUS, and BART—using four datasets: MultiNews, WikiSum, DUC, and CNN/Daily Mail. Their goal was to compare the abstractive summarization performance of these models on those datasets. According to their reported results, T5 achieved the strongest performance among the four models, with ROUGE-1, ROUGE-2, and ROUGE-L scores of 63.04%, 42.22%, and 47.82%, respectively, along with an average precision of 90.10%. The second-best performing model was BART, specifically the Facebook/BART-large-cnn variant.

Although the Transformer architecture has significantly advanced the field of natural language processing, the authors also identify several remaining challenges. They note that transformer-based summarization models may still fail to capture subtle contextual nuances and may produce summaries

Fig. 1: Proposed system architecture




Fig. 1: Proposed Architecture

that are less coherent than human-written ones. They also point out that fine-tuning can be computationally expensive and may not generalize equally well across different domains and text genres. In addition, evaluation metrics such as ROUGE may not fully reflect summary quality from a human perspective, which can create a gap between automatic evaluation scores and human judgment.

III. PROPOSED SYSTEM

This research proposes an automated meeting intelligence framework for processing meeting content from audio, video, and text inputs. The system combines Whisper-based transcription, hybrid summarization, structured information extrac-

tion, and domain classification to generate summaries, action items, deadlines, decisions, and domain labels across business, education, and healthcare-related settings.

Meeting audio is first transcribed into text using Whisper, after which the transcript undergoes preprocessing, summarization, and structured information extraction. This transcription stage is necessary because the summarization models operate on textual input rather than raw audio. Once the transcript is obtained, it is passed through a preprocessing stage and then into the summarization module.

In addition to abstractive generation, the system incorporates an extractive selection step to improve content coverage.

Using TF-IDF vectorization, each sentence in the original transcript is scored based on its cosine similarity to the overall document centroid. The top-ranked sentences are then combined with the abstractive output to produce a final hybrid summary that balances generated fluency with faithful coverage of key content from the source transcript.

The summarization stage employs transformer-based abstractive models, specifically FLAN-T5-base [23] and BART-samsum [24]. Unlike extractive methods, these models generate summaries in natural language rather than simply selecting sentences from the source transcript. This enables the system to produce shorter and more coherent summaries that better capture the main content of a meeting.

The proposed system is intended to support multiple application domains represented in this study, including business, education, and healthcare-related meeting content. Although the same summarization pipeline is applied across all domains, the inclusion of domain-representative data allows the system to be evaluated in different contextual settings.

The final output of the system includes a generated meeting summary together with extracted action items, deadlines, decisions, and predicted domain labels.

IV. DATASET

We used the AMI dataset from Hugging Face Datasets for training, validation, and testing [19]. The AMI Meeting Corpus [20], [21] is a multimodal dataset containing approximately 100 hours of English meeting recordings with aligned transcripts and summary-related annotations [20]. These properties make it suitable for meeting summarization research. However, a considerable portion of the dataset consists of scenario-based meetings, which may not fully represent all real-world settings [21].

Records with missing transcript or summary fields were excluded from the training and evaluation sets. After filtering, the final processed dataset consisted of 279 samples in total, divided into 209 training samples, 42 validation samples, and 28 test samples, formatted as input-target pairs for supervised fine-tuning.

In addition, the project included a smaller domain-specific collection of 40 meetings to represent the target application domains: 10 in education, 15 in business, and 15 in healthcare. It should be noted that the healthcare portion consisted of rough transcripts rather than original recorded meetings, since the corresponding audio and video were not available due to privacy concerns.

V. DATA PREPROCESSING

For audio inputs, the signal was loaded and resampled to 16 kHz using librosa [16], then segmented into 30-second chunks prior to transcription with Whisper small. For video inputs, the audio track was first extracted using FFmpeg and then processed using the same resampling and chunking procedure.

After transcription, several cleaning steps were applied to improve text quality before the transcripts were passed to the summarization models. First, common filler words were removed using regular expression matching. The targeted fillers included “um,” “uh,” “hmm,” “erm,” “ah,” “mm,” and “mhm,” which occur frequently in spontaneous speech but carry little semantic value for summarization. Second, repeated whitespace was collapsed into single spaces, and erroneous spaces before punctuation marks such as commas, periods, and question marks were corrected. Third, empty speaker lines, where a speaker tag appeared without any accompanying utterance, were removed to avoid introducing noise into the model input.

Following text cleaning, the transcripts were divided into smaller chunks to fit within the maximum input length supported by each model. Because BART-samsum supports inputs of up to 1,024 tokens, transcripts were segmented into chunks of approximately 750 words as a conservative heuristic before tokenization. For FLAN-T5-base, which supports up to 512 tokens, a smaller chunk size of approximately 400 words was used. These word-based chunk sizes were selected to reduce the likelihood of exceeding each model’s token limit during tokenization, particularly because meeting transcripts in the AMI corpus range from approximately 3,000 to over 20,000 words, far exceeding the input capacity of either model in a single pass. Each chunk was then independently tokenized and processed by the model, and the resulting chunk-level summaries were concatenated to form the full meeting summary.

VI. MODEL SELECTION

This project focused on abstractive summarization, in which the model generates a new summary rather than simply extracting sentences from the source text. For this reason, transformer-based sequence-to-sequence models were selected. Initially, larger variants of both BART and FLAN-T5 were considered. However, these models placed greater demands on the available notebook environment in terms of GPU memory and training time, and FLAN-T5-large also showed signs of overfitting in preliminary experiments. Because meeting transcripts frequently exceeded the maximum input length supported by the models, chunking was required, making computational efficiency and context capacity important practical considerations. Therefore, the final comparison was conducted using BART-samsum and FLAN-T5-base, which provided the most practical and stable configurations for this study. FLAN-T5-base was more computationally efficient, whereas BART-samsum offered a larger context window and stronger dialogue-oriented pretraining. This comparison was designed to examine the tradeoff between computational cost and summarization quality.

VII. MODEL TRAINING

The project was implemented in Python using Kaggle Notebook [12] with a T4 GPU, PyTorch [13], and Hugging Face Transformers [14]. The base working directory was defined to organize datasets, processed files, models, and evaluation

outputs. To improve reproducibility, random seeds were set for PyTorch and NumPy. tqdm [17] progress bars were used to monitor preprocessing, training, and evaluation progress within the notebook environment.

For the FLAN-T5 experiments, the google/flan-t5-base model was used. This model contains 247,577,856 parameters. The maximum input length was set to 512 tokens, and the maximum target summary length was set to 256 tokens. The model was trained for 20 epochs using a batch size of 4, gradient accumulation of 2, a learning rate of $5e-5$, and a weight decay of 0.01.

For the BART experiments, the philschmid/bart-large-cnn-samsum checkpoint was used. This model contains 406,290,432 parameters. The maximum input length was set to 1024 tokens, and the maximum target summary length was set to 512 tokens. The model was trained for 25 epochs using a batch size of 2, gradient accumulation of 4, a learning rate of $2e-5$, a weight decay of 0.01, a OneCycleLR learning-rate scheduler, and gradient clipping with a maximum norm of 1.0.

VIII. EVALUATION

The generated summaries were evaluated against reference summaries using ROUGE-1, ROUGE-2, and ROUGE-L [15]. These metrics are standard measures of summarization quality and estimate the degree of overlap between the generated summary and the human-written reference summary. In addition to automatic evaluation, sample outputs were also reviewed manually in order to assess coherence, relevance, and readability.

IX. PERFORMANCE


A. Metric Results

Table I: Model Comparison and Performance Results

Model	Parameters	ROUGE-1 Recall	ROUGE-2 Recall	ROUGE-L Recall	Max Input length
FLAN-T5-base	247.6M	71.22%	40.06%	43.32%	512
BART-samsum	406.3M	79.05%	45.64%	44.44%	1024

B. Model Comparison

Although BART-samsum contains more parameters and requires greater computational resources, it consistently outperformed FLAN-T5-base across all ROUGE metrics. The larger context window of BART-samsum and its pretraining on the SAMSum dialogue corpus likely contributed to its stronger performance on meeting transcripts.



C. Information Extraction Results

The pattern-based extraction module was applied to the full collection of 40 domain-specific meetings spanning business,

education, and healthcare. The module identified a total of 880 speaker-attributed action items, 15 key decisions, and 4 explicit deadline references throughout the collection. The higher volume of action items relative to decisions and deadlines reflects the natural distribution of these elements in meeting discourse. Meetings typically contain numerous task-oriented statements distributed across speakers, whereas formal decisions and explicitly stated deadlines occur less frequently in conversational dialogue. The AMI corpus, in particular, consists of design-team meetings where participants frequently express intentions and suggestions, which are appropriately captured by the action-item patterns.

X. RESULTS

The system was evaluated on the 28-sample test split of the AMI Meeting Corpus. Both BART-samsum and FLAN-T5-base were assessed using ROUGE-1, ROUGE-2, and ROUGE-L recall scores. BART-samsum consistently outperformed FLAN-T5-base across all three ROUGE metrics. On ROUGE-1 recall, BART-samsum achieved 79.05 percent compared to 71.22 percent for FLAN-T5-base, reflecting an improvement of approximately 7.83 percentage points. For ROUGE-2 recall, which captures bigram-level overlap and is more sensitive to phrasing quality, BART-samsum scored 45.64 percent versus 40.06 percent for FLAN-T5-base. The ROUGE-L recall scores were 44.44 percent and 43.32 percent respectively, indicating that both models captured comparable long-range subsequences from the reference summaries. BART-samsum also achieved higher F1 scores, with a ROUGE-1 F1 of 37.78 percent compared to 31.99 percent for FLAN-T5-base.

The higher recall performance of BART-samsum can be attributed to two factors. First, BART-samsum was pre-trained on the SAMSum dialogue corpus, which gives it an advantage in handling the turn-based, multi-speaker structure typical of meeting transcripts. Second, the hybrid summarization strategy adopted in this work, which combines abstractive chunk-level generation with TF-IDF-based extractive sentence selection, was designed to prioritize content coverage in order to minimize the risk of omitting important discussion points, decisions, or tasks from the final output.

Per-sample analysis further confirmed the robustness of BART-samsum, with ROUGE-1 recall scores ranging from 45.9 percent to 99.1 percent across the 28 test meetings. This variation is consistent with the diversity of meeting structures within the AMI corpus, including differences in meeting duration, number of speakers, and discussion complexity.

In addition to summarization, the pattern-based extraction module was applied to a collection of 40 domain-specific meetings spanning business, education, and healthcare. The module successfully identified a total of 880 speaker-attributed action items, 15 key decisions, and 4 explicit deadline references across the full collection. Each action item was linked to the corresponding speaker using speaker tags present in the transcript, enabling the generation of structured, assignee-

specific task lists. The keyword-based domain classifier correctly categorized all 40 meetings into their respective domains — business, education, or healthcare — based on term-frequency analysis.

The complete pipeline was deployed as a web-based interface using Gradio [18], supporting real-time input from video, audio, or text. The interface presents the generated summary, extracted action items with speaker assignments, detected deadlines, identified decisions, and the predicted domain label in a single unified view. Figures 2 and 3 illustrate the input and output interfaces of the deployed system.




Fig. 2: Input




Fig. 3: Final Result

XI. CONCLUSION

This paper presented an intelligent meeting summarization system for extracting structured information from audio, video, and text inputs. The main challenge addressed in this work is the development of a hybrid system that combines both abstractive and extractive methods. Audio is first converted into text using Whisper, an open-source speech recognition model developed by OpenAI. The resulting transcript is then processed using transformer-based summarization models, with both BART-samsum and FLAN-T5-base evaluated for their ability to generate coherent meeting summaries.

In addition to summarization, the system includes a pattern-matching module for extracting action items, deadlines, and key decisions from the transcript. Speaker information is used to associate tasks and deadlines with specific individuals, making the output more useful in practical settings. The system also applies keyword-based classification to identify whether a meeting belongs to the business, education, or healthcare domain. As demonstrated through the web-based interface, the complete pipeline can process meeting audio in real time.

While different from the sentence extraction approach outlined in the reference literature, which uses BERT as the backbone for extractive summarization, our approach is distinct because it generates new sentences rather than simply extracting existing ones from the transcript, resulting in more coherent and readable summaries. Our evaluation results on the AMI Meeting Corpus test dataset show that the BART-samsum model yielded a ROUGE-1 Recall score of 79.05 percent, which was higher than the score achieved by the FLAN-T5 baseline model and competitive with previously reported results in related summarization studies. Going forward, we intend to explore the incorporation of speaker diarization, the scaling up of training datasets through the use of larger and more diverse meeting datasets, and parameter-efficient fine-tuning.

REFERENCES

- [1] K. S. M. S and K. K K, "AI-based Meeting Notes Generator using Whisper and BERT," 2025 6th International Conference on IoT Based Control Networks and Intelligent Systems (ICICNIS), Bengaluru, India, 2025, pp. 1927-1931, doi: 10.1109/ICICNIS66685.2025.11315544.
- [2] P. V. Khandar, D. K. Mohanty, S. K. Baliarsingh, and V. V. Budati, "Exploring Abstractive Summarization of Pre-Trained Models: A Study on GPT-2, T5, Pegasus and BART," *International Journal of Applied Mathematics*, vol. 38, no. 1s, Oct. 2025, doi: 10.12732/ijam.v38i1s.749.
- [3] F. A. Darmawan, M. B. Mauludin, and C. S. K. Aditya, "Hybrid Video Transcription Summarization with a BERT-Based Clustering and BART," *Jurnal RESTI (Rekayasa Sistem dan Teknologi Informasi)*, vol. 9, no. 6, pp. 1439-1450, Dec. 2025, doi: 10.29207/resti.v9i6.7066.
- [4] A. S. Jadhav, A. B. Jha, R. R. Verma and P. Thakur, "A Comparative Framework for Educational Video Summarization using BART-CNN and PEGASUS," 2026 International Conference on AI-Driven Smart Systems and Ubiquitous Computing (ICAUC), Pathum Thani, Thailand, 2026, pp. 1825-1830, doi: 10.1109/ICAUC68182.2026.11441095.
- [5] Artificial Intelligence, "What is ROUGE? Meaning, Examples, Use Cases?," Feb. 20, 2026. [Online]. Available: <https://www.aiuniverse.xyz/rouge/>. [Accessed: Apr. 14, 2026].

- [6] Mindgrasp AI, “Mindgrasp – The #1 AI Study Tool for Students to Learn Faster,” Mindgrasp AI. [Online]. Available: mindgrasp.ai. [Accessed: Apr. 15, 2026].
- [7] Jumpshare, “Privacy Policy,” Nov. 13, 2024. [Online]. Available: <https://jumpshare.com/privacy>. [Accessed: Apr. 14, 2026].
- [8] Recal, “Privacy Policy,” Jan. 29, 2026. [Online]. Available: <https://tryrecal.com/privacy>. [Accessed: Apr. 14, 2026].
- [9] GoTranscript, “Audio & Video Summarizer (Free) — AI Summaries with Key Points.” [Online]. Available: <https://gotranscript.com/tools/audio-video-summarizer>. [Accessed: Apr. 13, 2026].
- [10] NoteGPT, “NoteGPT – AI Summarizer, Note Taking, PDF, PPT, YouTube, Video, Audio,” NoteGPT. [Online]. Available: notegpt.io. [Accessed: Apr. 15, 2026].
- [11] Read AI Help Center, “Security & Privacy Overview,” Mar. 18, 2026. [Online]. Available: <https://support.read.ai/hc/en-us/articles/25702259763091-Security-Privacy-Overview>. [Accessed: Apr. 13, 2026].
- [12] Kaggle, “Kaggle: The World’s AI Proving Ground,” Kaggle. [Online]. Available: <https://www.kaggle.com/>. [Accessed: Apr. 14, 2026].
- [13] PyTorch Contributors, “PyTorch,” PyTorch Documentation. [Online]. Available: <https://pytorch.org/>. [Accessed: Apr. 14, 2026].
- [14] Hugging Face, “Transformers,” Hugging Face Documentation. [Online]. Available: <https://huggingface.co/docs/transformers/index>. [Accessed: Apr. 14, 2026].
- [15] Hugging Face, “ROUGE,” evaluate GitHub Repository. [Online]. Available: <https://github.com/huggingface/evaluate/blob/main/metrics/rouge/README.md>. [Accessed: Apr. 14, 2026].
- [16] librosa development team, “librosa,” librosa Documentation. [Online]. Available: <https://librosa.org/index.html>. [Accessed: Apr. 14, 2026].
- [17] C. da Costa-Luis, “tqdm.rich,” tqdm Documentation. [Online]. Available: <https://tqdm.github.io/docs/rich/>. [Accessed: Apr. 14, 2026].
- [18] Gradio Team, “Build machine learning apps in Python,” Gradio. [Online]. Available: <https://www.gradio.app/>. [Accessed: Apr. 14, 2026].
- [19] Hugging Face, “edinburghcstr/ami,” [Online]. Available: <https://huggingface.co/datasets/edinburghcstr/ami>. [Accessed: Apr. 13, 2026].
- [20] AMI Project, “AMI Corpus,” Univ. of Edinburgh. [Online]. Available: <https://groups.inf.ed.ac.uk/ami/corpus/>. [Accessed: Apr. 14, 2026].
- [21] AMI Project, “AMI Corpus Overview,” Univ. of Edinburgh. [Online]. Available: <https://groups.inf.ed.ac.uk/ami/corpus/overview.shtml>. [Accessed: Apr. 14, 2026].
- [22] OpenAI, “Introducing Whisper,” Sep. 21, 2022. [Online]. Available: <https://openai.com/index/whisper/>. [Accessed: Apr. 15, 2026].
- [23] Hugging Face, “FLAN-T5,” Transformers Documentation. [Online]. Available: https://huggingface.co/docs/transformers/model_doc/flan-t5. [Accessed: Apr. 15, 2026].
- [24] Hugging Face, “bart-large-cnn-samsum,” Model Card for philschmid/bart-large-cnn-samsum. [Online]. Available: <https://huggingface.co/philschmid/bart-large-cnn-samsum>. [Accessed: Apr. 15, 2026].